SELF-SERVE COHORT · TAKE-HOME ANALYSIS

What the data is actually telling us.

PRESENTED BY

Tuomo Nikulainen

Head of Decision Intelligence · candidate

DATE

26 April 2026

THE DATA, IN THREE NUMBERS

There is no credit-leakage problem. You couldn't give the credits away if you tried.

CREDITS OFFERED

$2.0M

10,000 signups × $200

CREDITS CONSUMED

$52.1K

2.6% burn rate

SIGNUPS WITH ZERO API CALLS

81.6%

8,160 of 10,000

Source: self-serve cohort dataset · 6-month window ending 2026-02

THE QUESTION IS WRONG

"Are we attracting builders or freeloaders, and where is the credit spend going?"

THE TWO QUESTIONS WORTH ASKING

01 Why do four in five signups never touch the product?

THE ACTIVATION CLIFF

02 Why has conversion halved in five months while activation stayed flat?

THE FIRE ALARM

The rest of this deck is about those two — and what to build with the answer. Agents, not dashboards.

10,000 signups → 1,840 activations → 90 buyers → 20 SQLs.

The cliff sits at "signup → first API call."

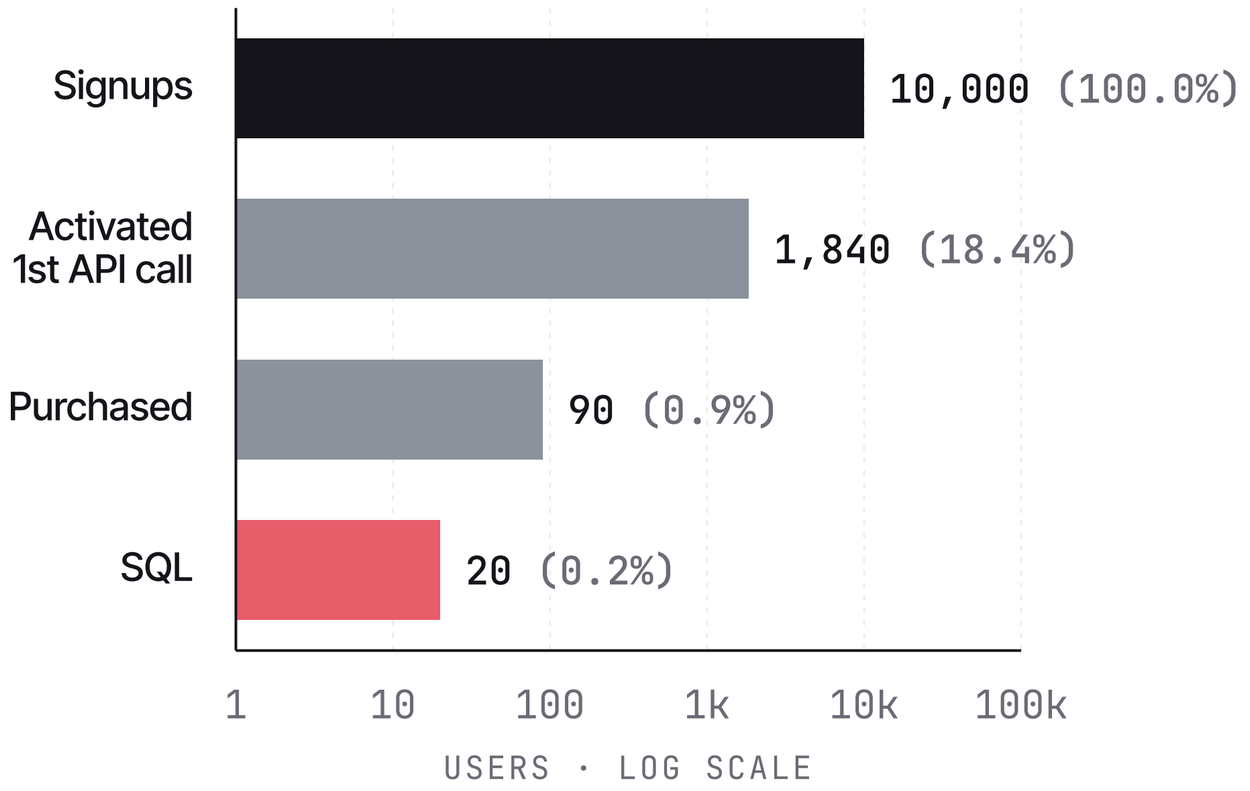

Top 1% of users consume 47% of all credits — and convert at 34%.

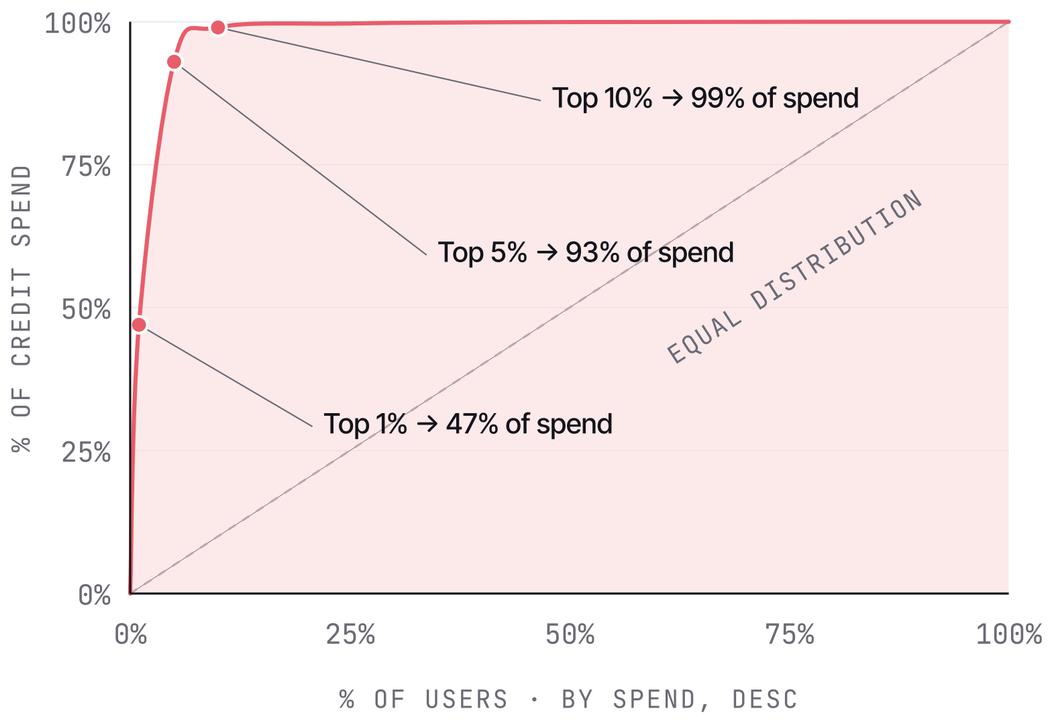

THE 1% THAT MATTERS

100 users · 47% of spend · 34% conversion · 100% on SDK or direct API.

The top 1% are the NRR base. The credit program already finds them — the leverage is finding more of them *earlier*, and expanding the ones we have.

Don't kill it. Compound it.

SECTION · 03

What predicts conversion — and what doesn't.

SDK adoption, voice_agent, business email, the broken SQL flag.

PREDICTORS OF CONVERSION

SDK adoption is the highest-leverage lever.

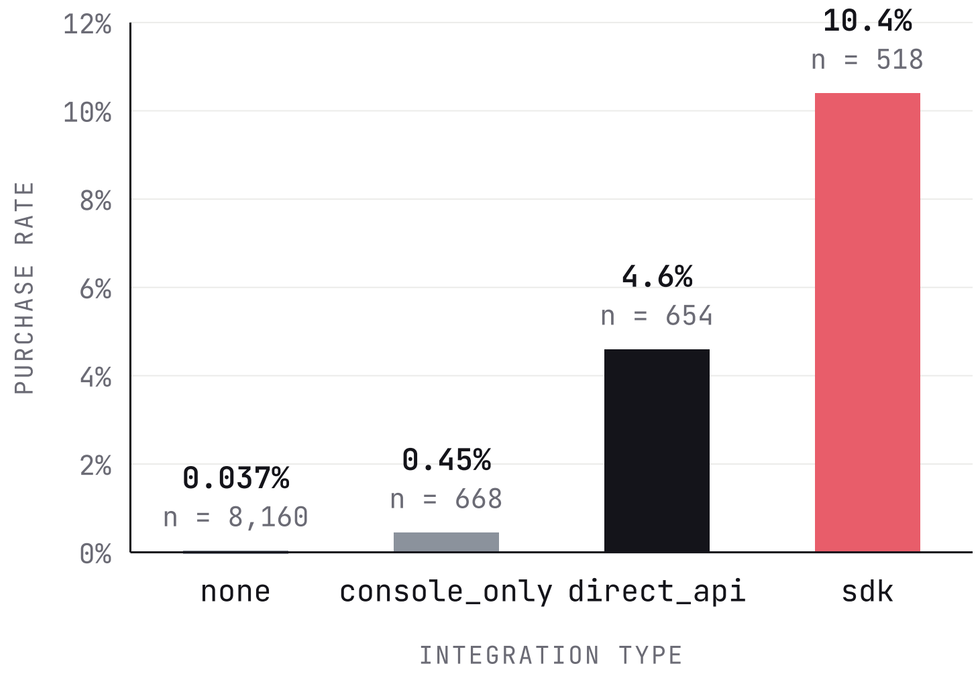

SIGNAL	CONVERSION LIFT	VERDICT
SDK adoption	282×	Killer feature
Direct API integration	125×	Strong
Business email domain	2.5×	Real but small
SIGNAL QUALITY — WHAT WE'RE NOT YET USING		
Current SQL flag	misses 99% of buyers	Definition is broken
`monthly_usage_pattern`	not yet segmented	Underused — tier-up agent
`voice_agent` product	5.8% conv. rate	Smallest cohort, hottest segment

Today only 5.2% of signups (518 / 10,000) ever install an SDK. Detailed `monthly_usage_pattern` segmentation pending notebook update.

COHORT-BY-COHORT, LAST 6 MONTHS

Conversion has halved since October. Activation hasn't moved.

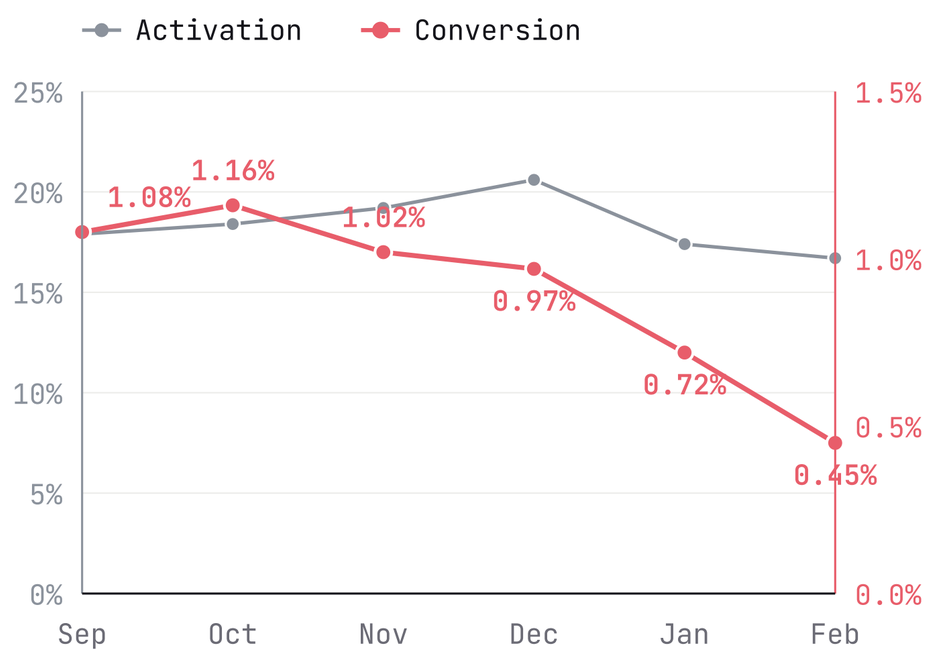

COHORT	ACTIVATION	CONVERSION
2025-09	17.9%	1.08%
2025-10	18.4%	1.16%
2025-11	19.2%	1.02%
2025-12	20.6%	0.97%
2026-01	17.4%	0.72%
2026-02	16.7%	0.45%

THREE HYPOTHESES · THE DATA ALONE CAN'T PICK

Something downstream of activation broke between October and February.

HYPOTHESIS 01

Pricing or paywall friction landed in Q1 2026.

Check release notes and the experiment log. Did a paywall, plan change, or rate-limit ship?

HYPOTHESIS 02

Onboarding regression broke the first-session "wow."

Same activation rate, fewer purchases — the gap is between first call and first dollar. Look at SDK quickstart and dashboard CTAs.

HYPOTHESIS 03

Competitive pressure compressed the consideration window.

OpenAI Realtime, ElevenLabs, Cartesia all moved during this window. Activated users may now be comparison-shopping.

Caveat handled: Sep had 6 months and converted 1.08%; Feb had 2 months and is at 0.45%. Median time-to-purchase is 4 days, so most Feb conversion has already happened. The drop is real, not censoring.

Next move: pull product release notes + experiment log; run an exit survey on the activated-not-converted Jan/Feb cohort.

53% of activators never come back for a second day. The ones who do convert at 27%.

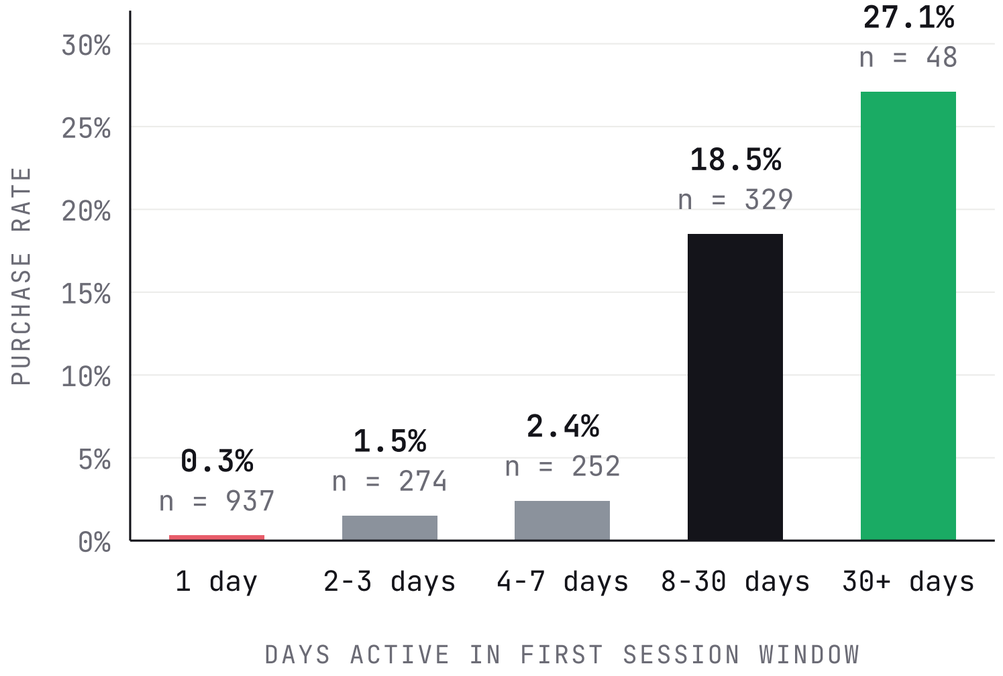

ONE DAY · THEN GONE

0.3%

934 users · they proved intent and disappeared

VS

30+ DAYS ACTIVE

27.1%

48 users · 90× the conversion rate per active day

THE LEVERAGE ZONE

Day-2 re-engagement is the obvious experiment. They came once. Bring them back.

RECOMMENDATION

Keep the $200 credit program. Redirect operational investment, in this order.

1 **Diagnose the conversion-rate decay first.**

Treat the Oct→Feb halving as an NRR-grade incident — not a marketing-funnel issue.

2 **Convert SDK adoption from a manual tutorial to a guided agent.**

282× conversion lift is sitting there; today only 5.2% of signups install one.

3 **Activate the latent NRR base.**

The 1,753 activated-not-purchased users (76% of all credit consumption) are expansion candidates, not freeloaders. Day-2 re-engagement is the first experiment.

4 **Replace the SQL definition with a continuously-learned scoring model.**

Current rules miss 99% of buyers. This is an automation bug, not a process gap.

One bet beyond the four: double down on `voice_agent` — highest conversion rate, smallest cohort, hottest market segment.

DASHBOARDS ARE A STEPPING STONE

Here's the destination. First three agents I'd ship.

01

Cohort Anomaly Sentinel

Watches activation, conversion, and spend by signup cohort.

TRIGGERS

a root-cause investigation when a cohort deviates >2σ from the trailing-12-week baseline.

CORRELATES

anomalies with deploy log, pricing changes, Salesforce signals, and — using Deepgram STT — sales-call transcripts.

LEARNS

from every confirmed cause; would have flagged the Oct→Feb decay in November.

02

Lead Scoring Loop

Continuously relearns the buyer signature from every new conversion.

REPLACES

the static SQL definition that today misses 99% of self-serve buyers.

OUTPUTS

a daily ranked list to AE and dev-rel queues.

CLOSES THE LOOP

— every accepted or rejected lead becomes training data.

03

Consumption Trajectory Agent

Reads `monthly_usage_pattern`-style data in real time.

IDENTIFIES

expansion candidates (growing MoM consumption) and churn risk (sudden drop).

TRIGGERS

in-product nudges, AE notifications, and pricing-tier offers — without human review for low-stakes actions.

DRIVES

NRR directly.

All three are deployable in the first 90 days against the existing data stack — **no infra rewrite required.**

THE ONE NUMBER THAT FLIPS THE RECOMMENDATION

Gross margin per customer.

ASSUMED ARR	REVENUE · 90 BUYERS	ROI ON $52K CREDITS
$500	$45K	0.9×
$2,500	$225K	4.3×
$5,000	$450K	8.6×
$25,000	$2.25M	43×

WHERE THE RECOMMENDATION FLIPS

From "hold the credit budget" to "expand it" — the answer sits between $2,500 and $5,000 ARR.

HOW I WORKED

I treated this analysis as an agentic-orchestration problem from the start: an explorer agent profiled the dataset, a critic agent stress-tested my framing against the prompt, and a builder agent drafted the artifacts under my review. I overrode the system on framing (the reframe in slide 03), prioritization, and the Intelligence Layer slide.

I run an AI agent failure-detection company in my day job — this is the operating mode I'd bring to Deepgram.

QUESTIONS I'D ASK YOU NEXT

What changed between Oct and Feb — pricing, onboarding, traffic mix?

What's the cost-to-serve on the 81.6% who never call the API?

Who currently owns the SDK-install moment, and what's their headcount?

Is the `voice_agent` 5.8% rate stable, or is the cohort still too small to trust?